

BackPropagation

There will be some functions that start with the word "grader" ex: grader_sigmoid(), grader_forwardprop(), grader_backprop() etc, you should not change those function definition.

Every Grader function has to return True.

Loading data

In [1]:

```
import pickle
import numpy as np
from tqdm import tqdm
import matplotlib.pyplot as plt
import math

with open('data.pkl', 'rb') as f:
    data = pickle.load(f)
print(data.shape)
X = data[:, :-1]
y = data[:, -1]
print(X.shape, y.shape)
```


(506, 6)
(506, 5) (506,)

In [30]:

```
A = data[0,:6]
print(A[:5],A[-1],'\n',A[:6])
```

[-1.2879095 -0.12001342 -1.45900038 -0.66660821 -0.14421743 1.858849127371369
[-1.2879095 -0.12001342 -1.45900038 -0.66660821 -0.14421743 1.85884913]

Computational graph



- If you observe the graph, we are having input features $[f_1, f_2, f_3, f_4, f_5]$ and 9 weights $[w_1, w_2, w_3, w_4, w_5, w_6, w_7, w_8, w_9]$.
- The final output of this graph is a value L which is computed as $(Y - Y')^2$

Task 1: Implementing backpropagation and Gradient checking

Check this video for better understanding of the computational graphs and back propagation

In [0]:


```
from IPython.display import YouTubeVideo
YouTubeVideo('i94OvYb6n0o',width="1000",height="500")
```


Out[0]:

- Write two functions
 - Forward propagation (Write your code in `def forward_propagation()`)


For easy debugging, we will break the computational graph into 3 parts.

Part 1






Part 2



Part 3



```

def forward_propagation(X, y, W):

    # X: input data point, note that in this assignment you are having 5-d data points
    # y: output varible
    # W: weight array, its of length 9, W[0] corresponds to w1 in graph, W[1] corresponds to w2 in gr
    # ... , W[8] corresponds to w9 in graph.
    # you have to return the following variables
    # exp= part1 (compute the forward propagation until exp and then store the values in exp)
    # tanh =part2(compute the forward propagation until tanh and then store the values in tanh)
    # sig =part3(compute the forward propagation until sigmoid and then store the values in sig)
    # now compute remaining values from computational graph and get y'
    # write code to compute the value of L=(y-y')^2
    # compute derivative of L w.r.t to Y' and store it in dl
    # Create a dictionary to store all the intermediate values
    # store L, exp,tanh,sig,dl variables

    return (dictionary, which you might need to use for back propagation)

```

- **Backward propagation(Write your code in def backward_propagation()) **

```

def backward_propagation(L, W,dictionary):

    # L: the loss we calculated for the current point
    # dictionary: the outputs of the forward_propagation() function

```

```

# dictionary, the outputs of the forward_propagation() function
# write code to compute the gradients of each weight [w1,w2,w3,...,w9]
# Hint: you can use dict type to store the required variables
# return dW, dW is a dictionary with gradients of all the weights

return dW

```

Gradient clipping

Check this [blog link](#) for more details on Gradient clipping

we know that the derivative of any function is

$$\lim_{\epsilon \rightarrow 0} \frac{f(x+\epsilon) - f(x-\epsilon)}{2\epsilon}$$

- The definition above can be used as a numerical approximation of the derivative. Taking an epsilon small enough, the calculated approximation will have an error in the range of epsilon squared.
- In other words, if epsilon is 0.001, the approximation will be off by 0.00001.

Therefore, we can use this to approximate the gradient, and in turn make sure that backpropagation is implemented properly. This forms the basis of gradient checking!

Gradient checking example

lets understand the concept with a simple example:

$$f(w1, w2, x1, x2) = w_1^2 \cdot x_1 + w_2 \cdot x_2$$

from the above function , lets assume $w_1 = 1, w_2 = 2, x_1 = 3, x_2 = 4$ the gradient of f w.r.t w_1 is

$$\begin{aligned} \frac{df}{dw_1} &= dw_1 \\ &= 2 \cdot w_1 \cdot x_1 \\ &= 2 \cdot 1 \cdot 3 \\ &= 6 \end{aligned}$$

let calculate the approximate gradient of w_1 as mentionned in the above formula and considering $\epsilon = 0.0001$

$$\begin{aligned} dw_1^{approx} &= \frac{f(w_1 + \epsilon, w_2, x_1, x_2) - f(w_1 - \epsilon, w_2, x_1, x_2)}{2\epsilon} \\ &= \frac{((1+0.0001)^2 \cdot 3 + 2 \cdot 4) - ((1-0.0001)^2 \cdot 3 + 2 \cdot 4)}{2\epsilon} \\ &= \frac{(1.00020001 \cdot 3 + 2 \cdot 4) - (0.99980001 \cdot 3 + 2 \cdot 4)}{2 \cdot 0.0001} \\ &= \frac{(11.00060003) - (10.99940003)}{0.0002} \\ &= 5.99999999999 \end{aligned}$$

$$\frac{\| (dW - dw^{approx}) \|_2}{\| (dW) \|_2 + \| (dw^{approx}) \|_2}$$

Then, we apply the following formula for gradient check: `gradient_check` =

The equation above is basically the Euclidean distance normalized by the sum of the norm of the vectors. We use normalization in case that one of the vectors is very small. As a value for epsilon, we usually opt for 1e-7. Therefore, if gradient check return a value less than 1e-7, then it means that backpropagation was implemented correctly. Otherwise, there is potentially a mistake in your implementation. If the value exceeds 1e-3, then you are sure that the code is not correct.

$$\text{in our example: } \text{gradient_check} = \frac{(6-5.99999999994898)}{(6+5.99999999994898)} = 4.2514140356330737e^{-13}$$

you can mathematically derive the same thing like this

$$\begin{aligned}
 dw_1^{approx} &= \frac{f(w_1 + \epsilon, w_2, x_1, x_2) - f(w_1 - \epsilon, w_2, x_1, x_2)}{2\epsilon} \\
 &= \frac{((w_1 + \epsilon)^2 \cdot x_1 + w_2 \cdot x_2) - ((w_1 - \epsilon)^2 \cdot x_1 + w_2 \cdot x_2)}{2\epsilon} \\
 &= \frac{4\epsilon \cdot w_1 \cdot x_1}{2\epsilon} \\
 &= 2 \cdot w_1 \cdot x_1
 \end{aligned}$$

Implement Gradient checking

(Write your code in `def gradient_checking()`)

Algorithm

`W = initialize_randomly`

`def gradient_checking(data_point, W):`

```

# compute the L value using forward_propagation()
# compute the gradients of W using backword_propagation()
approx_gradients = []
for each wi weight value in W:
    # add a small value to weight wi, and then find the values of L with the updated weights
    # subtract a small value to weight wi, and then find the values of L with the updated weights
    # compute the approximation gradients of weight wi
    approx_gradients.append(approximation gradients of weight wi)
# compare the gradient of weights W from backword_propagation() with the aproximation gradients of weights with
<br> gradient_check formula
return gradient_check

```

NOTE: you can do sanity check by checking all the return values of `gradient_checking()`, they have to be zero. if not you have bug in your code

Task 2 : Optimizers

- As a part of this task, you will be implementing 3 type of optimizers(methods to update weight)
- Use the same computational graph that was mentioned above to do this task
- Initilize the 9 weights from normal distribution with mean=0 and std=0.01

Check below video and [this blog](#)

In [15]:

```
from IPython.display import YouTubeVideo
YouTubeVideo('gYpoJMggyXA',width="1000",height="500")
```

Out[15]:



Algorithm

```
for each epoch(1-100):
    for each data point in your data:
        using the functions forward_propagation() and backward_propagation() compute the gradients of weights
        update the weights with help of gradients ex: w1 = w1 - learning_rate * dw1
```

Implement below tasks

- Task 2.1: you will be implementing the above algorithm with Vanilla update of weights
- Task 2.2: you will be implementing the above algorithm with Momentum update of weights
- Task 2.3: you will be implementing the above algorithm with Adam update of weights

Note : If you get any assertion error while running grader functions, please print the variables in grader functions and check which variable is returning False .Recheck your logic for that variable .

Task 1

Forward propagation

In [4]:

```
def sigmoid(z):
    """In this function, we will compute the sigmoid(z)"""
    # we can use this function in forward and backward propagation

    return 1/(1+math.exp(-z))
```

In [5]:

```
def forward_propagation(x, y, w):
    """In this function, we will compute the forward propagation """
    # X: input data point, note that in this assignment you are having 5-d data points
    # y: output variable
    # W: weight array, its of length 9, W[0] corresponds to w1 in graph, W[1] corresponds to w2 in graph,..., W[8] corresponds to w9 in graph.
```

```

# you have to return the following variables
# exp=part1 (compute the forward propagation until exp and then store the values in exp)
# tanh =part2(compute the forward propagation until tanh and then store the values in tanh)
# sig =part3(compute the forward propagation until sigmoid and then store the values in sig)
# now compute remaining values from computational graph and get y'
# write code to compute the value of L=(y-y')^2
# compute derivative of L w.r.to Y' and store it in dl
# Create a dictionary to store all the intermediate values
# store L, exp,tanh,sig variables
#return (dictionary, which you might need to use for back propagation)
q3 = (w[0]*x[0])+(w[1]*x[1])
q6 = (q3**2) + w[5]
if q6>700:
    q8 = math.exp(700)+w[6]
else:
    q8 = math.exp((q3**2)+w[5])+w[6]
q20 = w[2]*x[2]
q19 = math.sin(q20)
q16 = (w[3]*x[3])+(w[4]*x[4])
q14 = (q16*q19)+w[7]
q13 = sigmoid(q14)

temp = ((w[0]*x[0]+w[1]*x[1])**2)+w[5]
if temp>700:
    temp = 700

exp = math.exp(temp)
tanh = math.tanh(exp+w[6])
sig = sigmoid(math.sin(w[2]*x[2])*((w[3]*x[3])+(w[4]*x[4]))+w[7])
y_hat = tanh+(sig*w[8])
L = (y-y_hat)**2
dl = 2*(y_hat-y)

interm_values = {'dl':dl,'loss':L,'exp':exp,'tanh':tanh,'sigmoid':sig,'q3': q3,'q6': q6,'q8': q8,'q20': q20,
                 'q19': q19,'q16': q16,'q14': q14,'q13': q13}
return interm_values

```

Grader function - 1

In [6]:

```

def grader_sigmoid(z):
    val=sigmoid(z)
    assert(val==0.8807970779778823)
    return True
grader_sigmoid(2)

```

Out[6]:

True

Grader function - 2

In [7]:

```

def grader_forwardprop(data):
    dl = (np.round(data['dl'],4)==-1.9285)
    loss=(np.round(data['loss'],4)==0.9298)
    part1=(np.round(data['exp'],4)==1.1273)
    part2=(np.round(data['tanh'],4)==0.8418)
    part3=(np.round(data['sigmoid'],4)==0.5279)
    assert(dl and loss and part1 and part2 and part3)
    return True
w=np.ones(9)*0.1
dl=forward_propagation(X[0],y[0],w)
grader_forwardprop(dl)

```

Out[7]:

True

Backward propagation

In [8]:

```
def backward_propagation(L,w,dict):
    """In this function, we will compute the backward propagation """
    # L: the loss we calculated for the current point
    # dictionary: the outputs of the forward_propagation() function
    # write code to compute the gradients of each weight [w1,w2,w3,...,w9]
    # Hint: you can use dict type to store the required variables
    # dw1 = # in dw1 compute derivative of L w.r.to w1
    # dw2 = # in dw2 compute derivative of L w.r.to w2
    # dw3 = # in dw3 compute derivative of L w.r.to w3
    # dw4 = # in dw4 compute derivative of L w.r.to w4
    # dw5 = # in dw5 compute derivative of L w.r.to w5
    # dw6 = # in dw6 compute derivative of L w.r.to w6
    # dw7 = # in dw7 compute derivative of L w.r.to w7
    # dw8 = # in dw8 compute derivative of L w.r.to w8
    # dw9 = # in dw9 compute derivative of L w.r.to w9
    temp = dict['q6']
    if temp>700:
        temp = 700
    dw1 = (2*dict['q3']*L[0])*((math.tanh(dict['q8'])**2)-1)*math.exp(temp)*(-1*dict['dl'])
    dw2 = (2*dict['q3']*L[1])*((math.tanh(dict['q8'])**2)-1)*math.exp(temp)*(-1*dict['dl'])
    dw3 = dict['dl']*w[8]*(sigmoid(dict['q14'])*(1-sigmoid(dict['q14']))) *dict['q16']*math.cos(dict['q20'])*L[2]
    dw4 = dict['dl']*w[8]*(sigmoid(dict['q14'])*(1-sigmoid(dict['q14']))) *dict['q19']*L[3]
    dw5 = dict['dl']*w[8]*(sigmoid(dict['q14'])*(1-sigmoid(dict['q14']))) *dict['q19']*L[4]
    dw6 = dict['dl']*(1-(math.tanh(dict['q8'])**2))*math.exp(temp)
    dw7 = dict['dl']*(1-(math.tanh(dict['q8'])**2))
    dw8 = dict['dl']*w[8]*(sigmoid(dict['q14'])*(1-sigmoid(dict['q14'])))
    dw9 = dict['dl']*dict['q13']

    dW = {'dw1':dw1,'dw2':dw2,'dw3':dw3,'dw4':dw4,'dw5':dw5,'dw6':dw6,'dw7':dw7,'dw8':dw8,'dw9':dw9}
    return dW
```

Grader function - 3

In [9]:

```
def grader_backprop(data):
    dw1=(np.round(data['dw1'],8)==-0.22973323)
    dw2=(np.round(data['dw2'],8)==-0.02140761)
    dw3=(np.round(data['dw3'],8)==-0.00562541)
    dw4=(np.round(data['dw4'],8)==-0.00465794)
    dw5=(np.round(data['dw5'],8)==-0.00100772)
    dw6=(np.round(data['dw6'],8)==-0.63347519)
    dw7=(np.round(data['dw7'],8)==-0.56194184)
    dw8=(np.round(data['dw8'],8)==-0.04806288)
    dw9=(np.round(data['dw9'],8)==-1.01810444)
    assert(dw1 and dw2 and dw3 and dw4 and dw5 and dw6 and dw7 and dw8 and dw9)
    return True
w=np.ones(9)*0.1
d1=forward_propagation(X[0],y[0],w)
d1=backward_propagation(X[0],w,d1)
grader_backprop(d1)
```

Out[9]:

True

Implement gradient checking

In [48]:

```
def gradient_checking(data_point, w):
    # compute the L value using forward_propagation()
    # compute the gradients of W using backward_propagation()
    x = data_point[:5]
    ...
```

```

y = data_point[-1]
d1 = forward_propagation(x,y,w)
d2 = backward_propagation(x,w,d1)
W = np.array([d2[i] for i in d2.keys()])

approx_gradients = []
epsilon = 1e-7
for i in range(len(w)):
    # add a small value to weight wi, and then find the values of L with the updated weights
    # subtract a small value to weight wi, and then find the values of L with the updated weights
    # compute the approximation gradients of weight wi
    w[i] = w[i]+epsilon
    d3 = forward_propagation(x,y,w)
    loss = d3['loss']

    w[i] = w[i]-((epsilon)*2)
    d4 = forward_propagation(x,y,w)
    loss2 = d4['loss']
    w[i]+=epsilon

    to_app = (loss-loss2)/(2*epsilon)
    approx_gradients.append(to_app)

# compare the gradient of weights W from backward_propagation() with the approximation gradients of weights with gradient_check formula
la

gradient_checking = []
for i in range(len(W)):
    num = W[i]-approx_gradients[i]
    den = W[i]+approx_gradients[i]
    check = num/den
    if check < 1e-7:
        print('True')
    gradient_checking.append(check)

return gradient_checking

```

In [49]:

```

import random
w = 0.01*np.random.randn(9)+0
gradient_checking(data[0],w)

```

```

True
True
True
True
True
True
True
True
True

```

Out[49]:

```

[1.740006799697927e-08,
 -3.6508454843836556e-07,
 -3.41869346193698e-06,
 -1.1193186678764233e-05,
 -1.1095160699052884e-05,
 -1.6901170504871824e-10,
 2.1193841152009169e-10,
 -7.147585745262382e-08,
 -7.591098230085461e-11]

```

Task 2: Optimizers

for each epoch(1-100): for each data point in your data: using the functions forward_propagation() and backward_propagation() compute the gradients of weights update the weights with help of gradients ex: $w_1 = w_1 - \text{learning_rate} * dw_1$

Algorithm with Vanilla update of weights

Algorithm with Vanilla update of weights


In [44]:

```
#Vanilla update
import random
w = 0.01*np.random.randn(9)+0
epochs = 100
l_rate = 0.0001
Loss1 = []
for i in range(epochs):
    for j,k in enumerate(data):
        X = data[j][:-5]
        y = data[j][-1]
        d1=forward_propagation(X,y,w)
        d2=backward_propagation(X,w,d1)
        dw = np.array([d2[i] for i in d2.keys()])
        w = w -(l_rate*dw)
    Loss1.append(d1['loss'])
```

Plot between epochs and loss

In [45]:

```
plt.title('Vanilla update')
plt.xlabel('epochs')
plt.ylabel('loss')
plt.plot(range(epochs),Loss1)
plt.show()
```



Algorithm with momentum update of weights


In [46]:

```
import random
w = 0.01*np.random.randn(9)+0
epochs = 100
l_rate = 0.0001
gamma = 0.9
Loss2 = []
momentum = 0
for i in range(epochs):
    for j,k in enumerate(data):
        X = data[j][:-5]
        y = data[j][-1]
        d1=forward_propagation(X,y,w)
        d2=backward_propagation(X,w,d1)
        dw = np.array([d2[i] for i in d2.keys()])
        dw = l_rate*dw
        w = w -(momentum+dw)
        momentum = gamma*dw
    Loss2.append(d1['loss'])
```

Plot between epochs and loss

In [47]:

```
plt.title('Momentum update')
plt.xlabel('epochs')
plt.ylabel('loss')
plt.plot(range(epochs),Loss2)
plt.show()
```



Algorithm with Adam update of weights


In [48]:


```
import random
w=0.01*np.random.randn(9)+0
epochs = 100
l_rate = 0.0001
B1 = 0.9
B2 = 0.99
mt = 0
vt = 0
count = 0
Loss3 = []
for i in range(epochs):
    for j,k in enumerate(data):
        count+=1
        X = data[j][:5]
        y = data[j][-1]
        d1=forward_propagation(X,y,w)
        d2=backward_propagation(X,w,d1)
        dw = np.array([d2[i] for i in d2.keys()])
        mt = (B1*mt)+(1-B1)*dw
        vt = (B2*vt)+(1-B2)*(dw**2)
        mt_hat = mt/(1-(B1**count))
        vt_hat = vt/(1-(B2**count))
        w = w - (l_rate*(mt_hat/(vt_hat**0.5)+1e-8))
    Loss3.append(d1['loss'])
```

Plot between epochs and loss

In [49]:

```
plt.title('Adam update')
plt.xlabel('epochs')
plt.ylabel('loss')
plt.plot(range(epochs),Loss3)
plt.show()
```





Comparision plot between epochs and loss with different optimizers

In [55]:

```
plt.title('Comparision plot of Optimizers')
plt.plot(range(epochs),Loss1,'b',label='Vanilla')
plt.plot(range(epochs),Loss2,'g',label='Momentum')
plt.plot(range(epochs),Loss3,'r',label='Adam')
plt.legend()
plt.show()
```

